

DEEP LEARNING CNN


RESEARCHGROUP
I PRODAM3D

Cristian López Del Alamo
clopezd@utec.edu.pe

Noviembre, 2021

Schedule

1. Introduction
2. Kernel
3. Convolution
4. Pooling
5. CNN
6. Architecture
7. Applications



Fuente: [Click](#)

Faces with different emotions



[Fuente: Click](#)

1. Introduction



Image



0.1	0.8	1.3	.3	3.1	4.2
-----	-----	-----	----	-----	-----

Feature vector

Wavelets, Harris, SIFT,
Histogram color, etc



Wavelets, Harris, SIFT,
Histogram color, etc

0.1	0.8	1.3	.3	3.1	4.2
-----	-----	-----	----	-----	-----

v1: Feature Vector



0.2	0.9	1.1	.2	3.5	4.0
-----	-----	-----	----	-----	-----

v2: Feature Vector

$$d(v1, v2) < e$$

Wavelets, Harris, SIFT,
Histogram color, etc



0.1	0.8	1.3	.3	3.1	4.2
-----	-----	-----	----	-----	-----

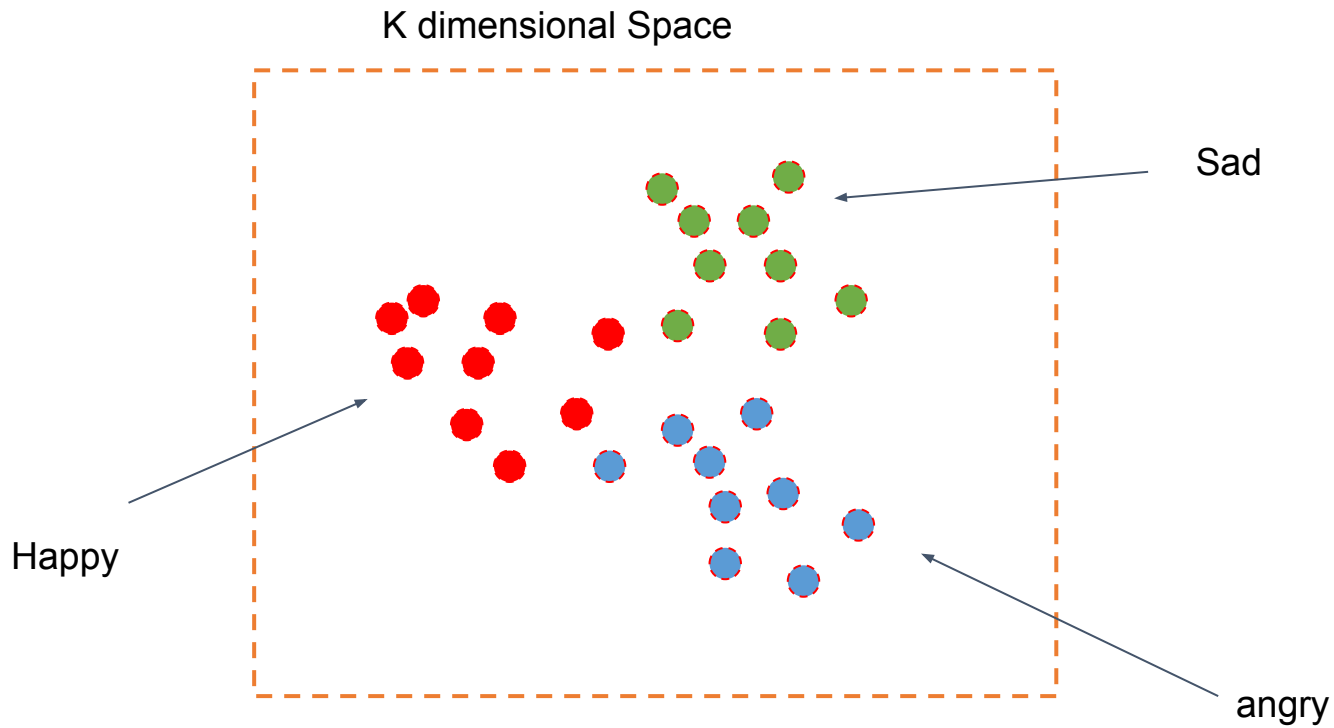
v1: Feature Vector




2.2	3.9	1.6	2.2	6.5	1.0
-----	-----	-----	-----	-----	-----

v2: Feature Vector

$d(v1, v2) > e$





C. López, L. Arnaldo, L. Fuentes, W. Ramos, "Agrupamiento por similitud de imágenes mediante Árbol de Expansión Mínima y Soft Heap", XLI Latin American Computing Conference, pp. 1-7, 2013



X



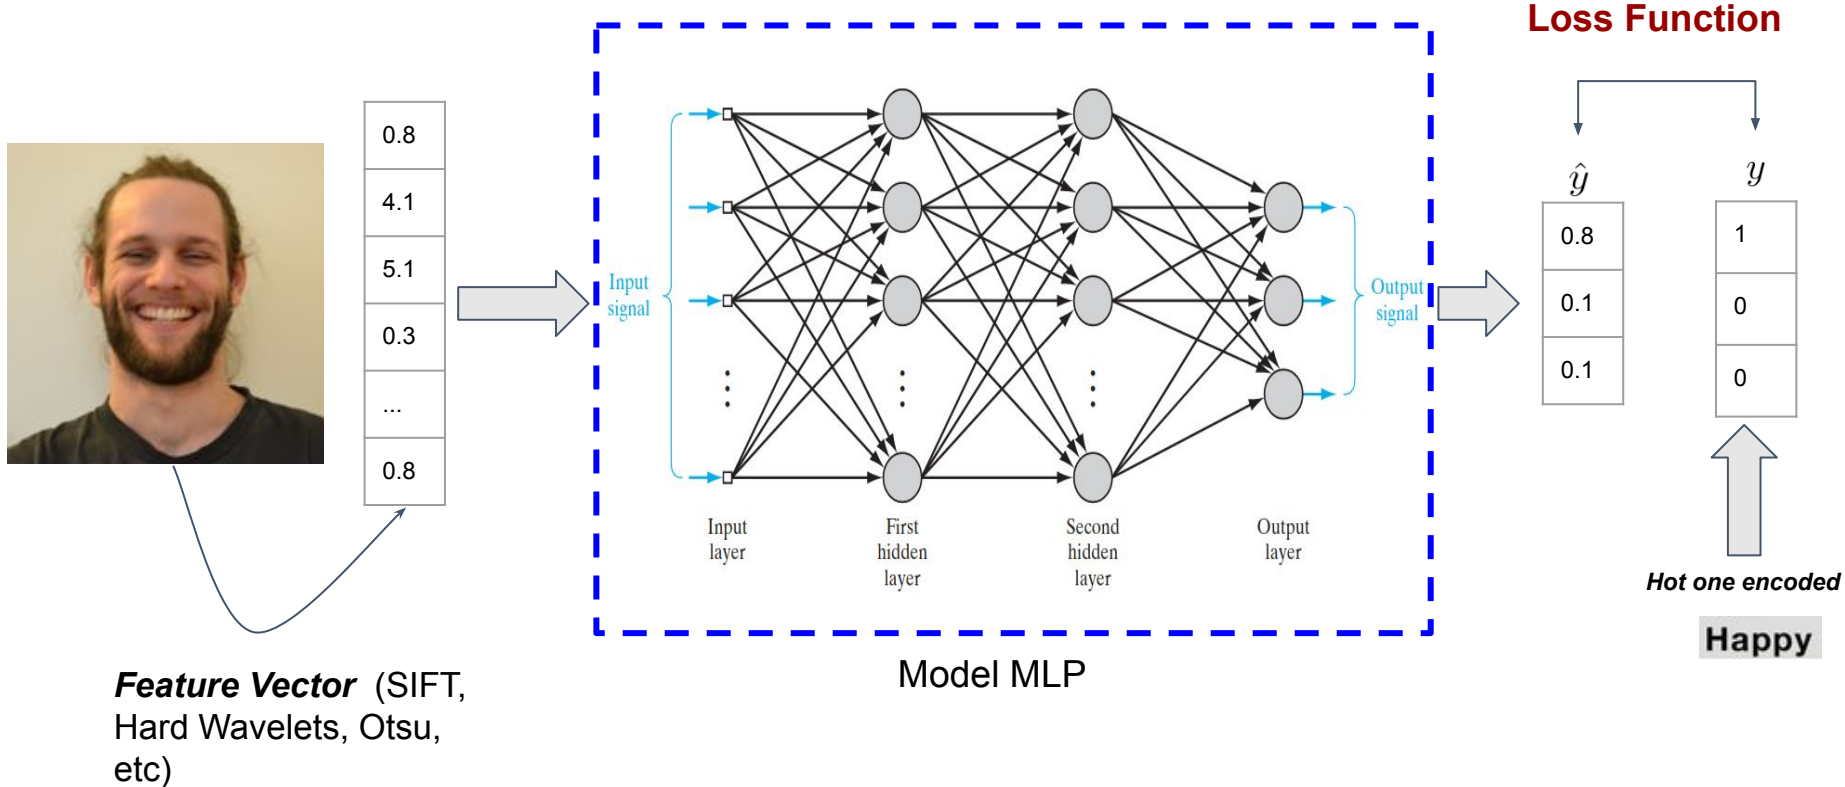
Neuronal Network



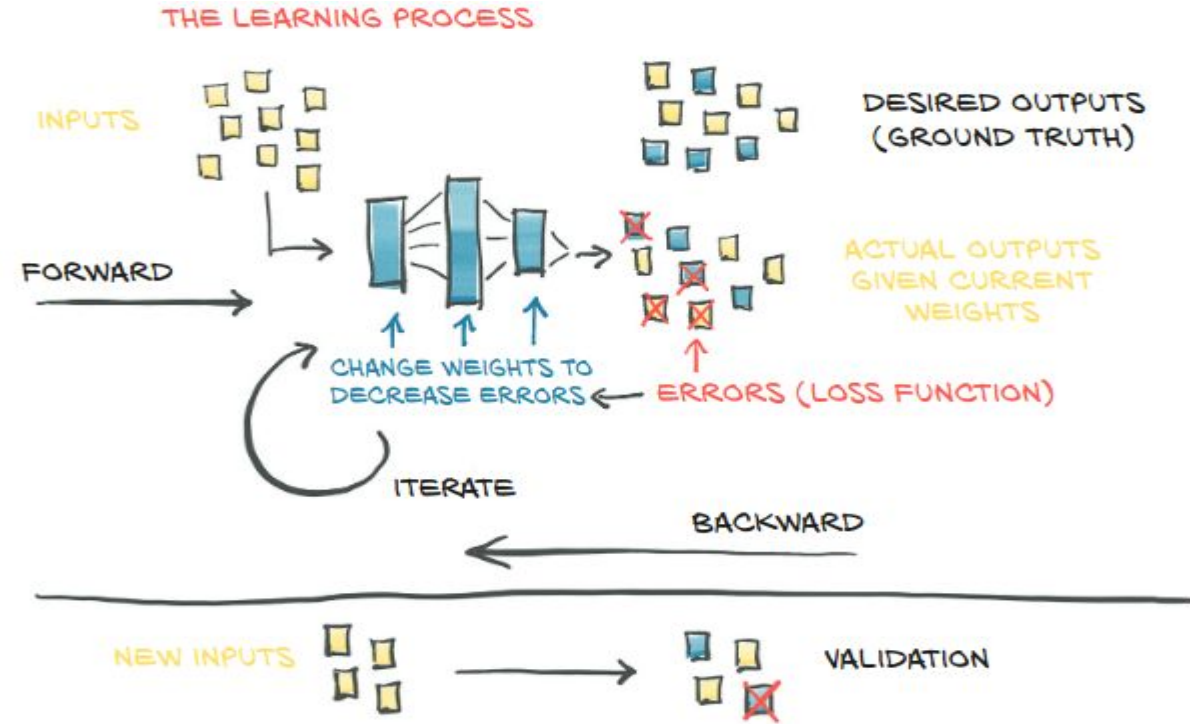
Happy

Y

Multilayer Perceptron




Multilayer Perceptron




Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Multilayer Perceptron




[Fuente: Eli Stevens, et al, Deep Learning with PyTorch](#)

Multilayer Perceptron



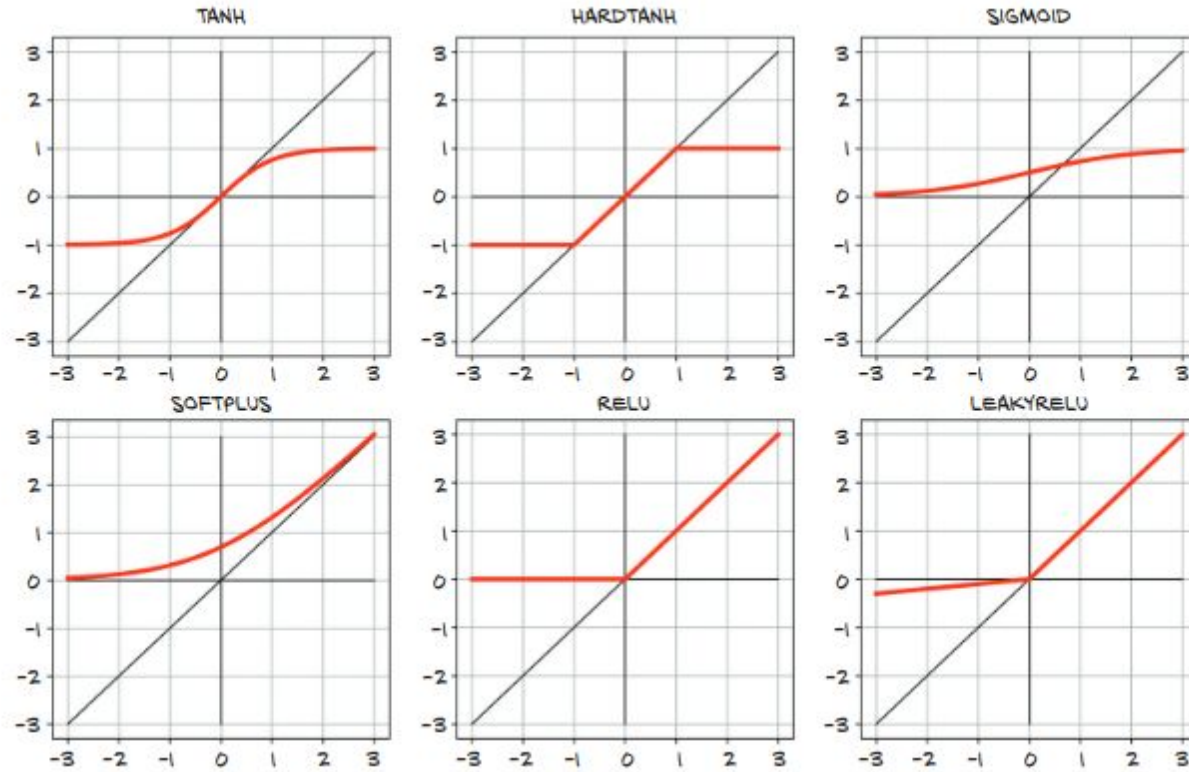
Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Activation Function



Fuente: [Eli Stevens, et al., Deep Learning with PyTorch](#)

Activation Function




[Fuente: Eli Stevens, et all, Deep Learning with PyTorch](#)

Problems MLP



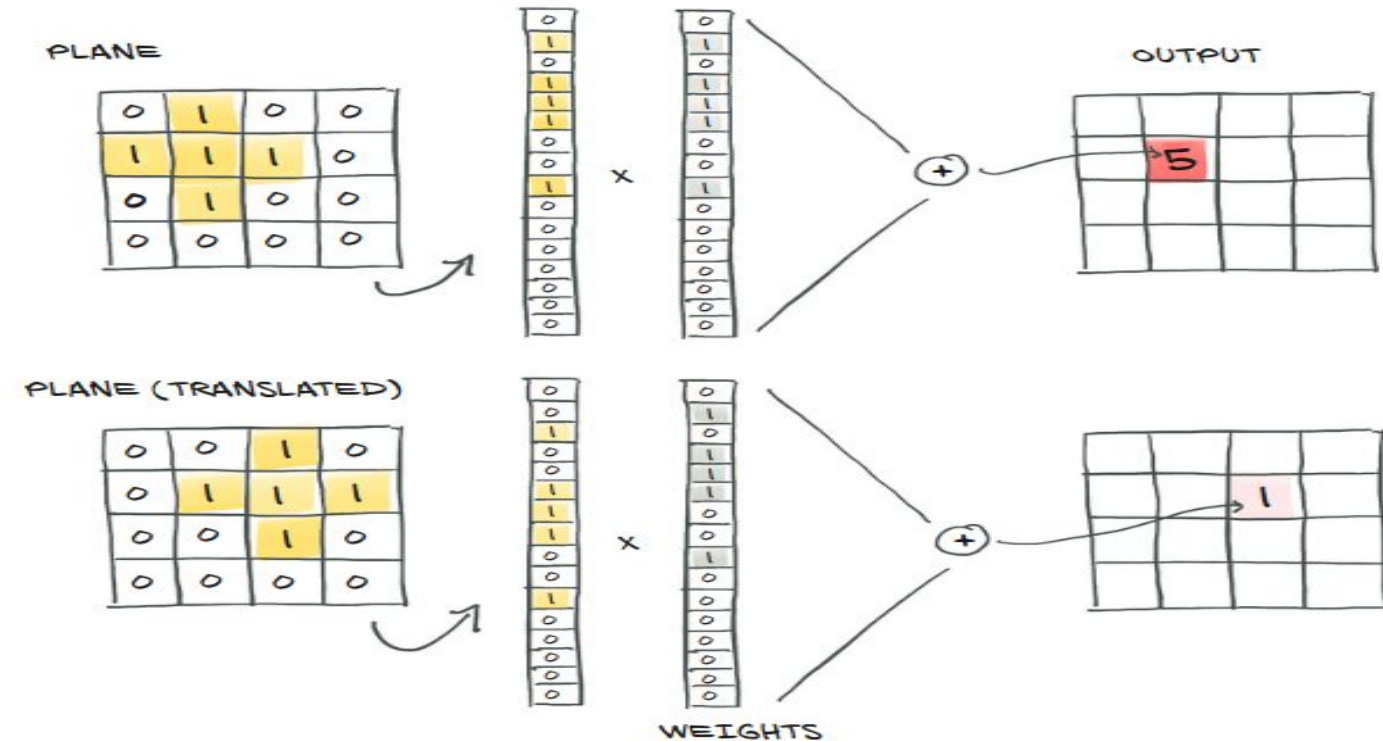
[Fuente: Click](#)

Problems with MLP




Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Problems with MLP



Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Problems with MLP



Ian Goodfellow, Yoshua Bengio, Aaron Courville; Deep Learning, 2016

CONVOLUTIONAL NEURAL NETWORK CNN


RESEARCHGROUP
I PRODAM3D

Dr. Cristian López Del Alamo

January 4, 2021



2. Kernel

Kernel





0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0

(a)



0	1	2	1	0
1	3	5	3	1
2	5	9	5	2
1	3	5	3	1
0	1	2	1	0

(b)



0	0	-1	0	0
0	-1	-2	-1	0
-1	-2	16	-2	-1
0	-1	-2	-1	0
0	0	-1	0	0

(c)

Example

3. Convolution

Convolution

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

0	-1	0
-1	5	-1
0	-1	0

114				

[Fuente: Click](#)

Convolution

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

↓
308

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

↓
-498

+

0	1	1
0	1	0
1	-1	1

Kernel Channel #3


↓
164
↑
Bias = 1
+ 1 = -25

Output

-25				...
				...
				...
				...
...

[Fuente: Click](#)

Convolution



Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Convolution layer : Stride




Image size = (7x7)
Filter size = (3x3)
padding = 0
stride = 1

Convolution layer : Stride




Image size = (7x7)
Filter size = (3x3)
padding = 0
stride = 1

Convolution layer : Stride




Image size = (7x7)
Filter size = (3x3)
padding = 0
stride = 1

Convolution layer : Stride




Image size = (7x7)
Filter size = (3x3)
padding = 0
stride = 1

Convolution layer : Stride




Image size = (7x7)

Filter size = (3x3)

padding = 0

stride = 1

output size = ?

Convolution layer : Stride




Image size = (7x7)
Filter size = (3x3)
padding = 0
stride = 1

output size = 5x5

Convolution layer : Stride




Image size = (7x7)

Filter size = (3x3)

padding = 0

stride = 2

Convolution layer : Stride




Image size = (7x7)
Filter size = (3x3)
padding = 0
stride = 2

Convolution layer : Stride





Image size = (7×7)
Filter size = (3×3)
padding = 0
stride = 2

output size = 3×3

What happen if the kernel size is 3×3 ?

Convolution layer : Stride



Convolution layer : Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Image size = (7x7)

Filter size = (3x3)

padding = 1

stride = 1

What the output size?

4. Pooling

Pooling

INPUT
(OUTPUT OF CONV + ACTIVATION)

2	2	2	0
2	5	2	1
2	2	2	0
0	1	0	0

MAXPOOL



OUTPUT

2	2
2	5

MAX=5

MAX=2

2	2
0	1

2	0
2	1

MAX=2

MAX=2

2	0
0	0

5	2
2	2

Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Pooling layer





image source: [click](#)

- Invariant to small translation of the input
- Can handle inputs of variable size and return outputs of the same size

5. CNN

Convolution and Pooling



Fuente: Eli Stevens, et all, [Deep Learning with PyTorch](#)

Typical layer in CNN





image source: [click](#)




image source: [click](#)


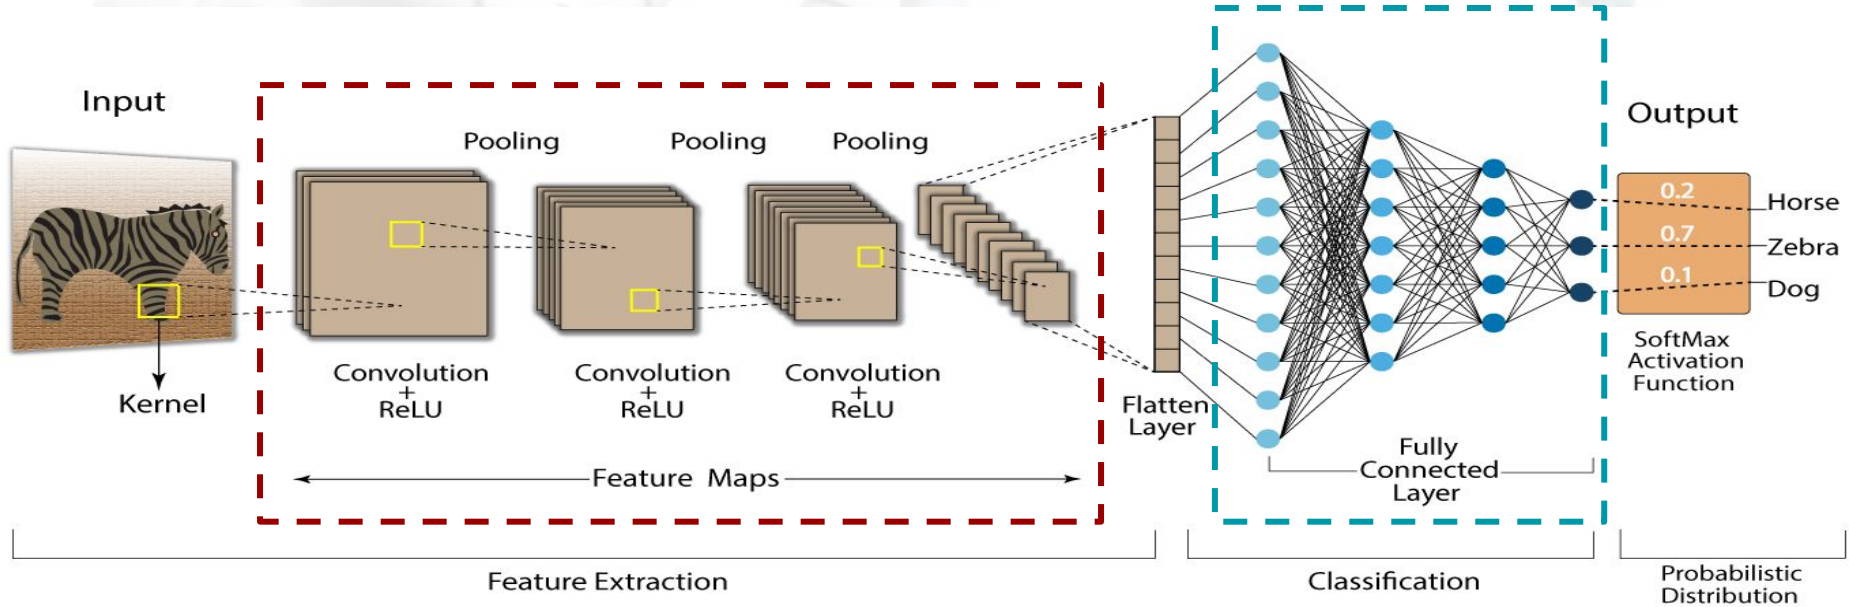


image source: [click](#)

6. Architecture

Architecture



$$f\left(\text{Input Image}\right) = \hat{y} \quad |y - \hat{y}|_p^p < \epsilon$$

CNN Architecture


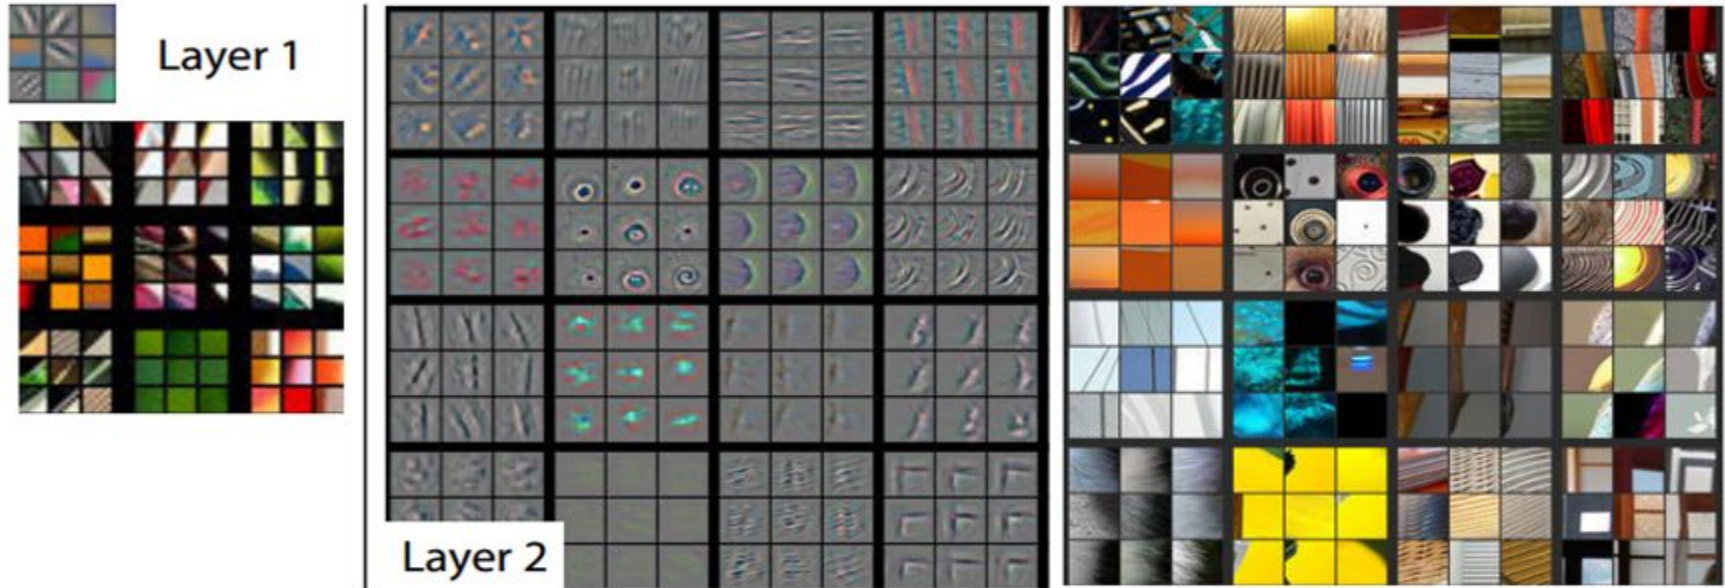


image source: [click](#)

CNN Architecture



Visualizations of Layer 1 and 2. Each layer illustrates 2 pictures, one which shows the filters themselves and one that shows what part of the image are most strongly activated by the given filter. For example, in the space labeled Layer 2, we have representations of the 16 different filters (on the left)

image source: [click](#)

7. Applications

Applications





Image Segmentation




Object Detection and Segmentation

Applications



Image Segmentation



Video Tracking

Applications




Denoising




Super resolution

Applications



Inpainting

INTRODUCTION




ImageNet Large Scale Visual Recognition Challenge results show that deep learning is surpassing human levels of accuracy.

Source: [click](#)

Pythor Example

CNN Pythorh



[Colab Example](#)

Extra Information

[CNN Backpropagation derivatives](#)